

The early signals are promising.

Selected highlights from our
API and Codex evaluations at Eliza.


364 candidate attempts · 18.4M recorded tokens

Early access, real experiments · Teja

Input + output, including cached and repeated context, across both models. Tested Vega-alpha, now named GPT-6 Astra. Selected exploratory findings; results vary by task.

Exact reasoning.

64.2% fewer output tokens.





Same acceptance checks. **Both models passed every run in this slice.**

Matched passing pairs; change in summed output tokens, including reasoning. Three designs with two repetitions each. Same configured effort does not imply equal compute.

Efficiency signals in Codex, too.

Public comparison scope: 322 API attempts · 42 Codex assignments


WITHIN CODEX · OUTPUT TOKENS INCLUDING REASONING · VEGA VS SOL



Codex used synthetic packet tasks with unequal runtime inputs. These selected high-effort comparisons are separate from API results; different tasks and context prevent a causal comparison between surfaces.

Long context.

29.0% fewer output tokens.



Both models passed across the two context sizes.

High vs high · 2 designs at nominal 96k and 192k context sizes · 4 passing pairs. This comparison is output only, including reasoning; input and cached context are excluded.

A more natural delegation workflow.

Eliza teammates described Astra using
subagents more naturally and needing
less steering on longer implementations.

Plan → Delegate → Coordinate

A practitioner signal worth testing next.

Qualitative reports from Eliza engineers, not a measured benchmark result. The candidate harness assigned scoped tasks; it did not measure autonomous subagent usage or coordination.

Match the model to the task.

Exact constraint reasoning

3 designs · 6 passing pairs

Earlier technical API packets

14 designs · 14 passing pairs

Public coding subset

16 designs · 16 passing pairs

Staged repository work

3 clean workflows · 3 passing pairs

Fewer Vega tokens

More Vega tokens



Same effort within each row: **high vs high.**

Matched passing pairs; summed output tokens, including reasoning. These are separate task families, not a pooled effect. Efficiency gains vary by workload; repeats are not new designs.

The evidence should travel with the claim.

01

Task

Save the prompt.
Define acceptance.
Record effort and limits.

02

Run

Keep provider usage.
Separate API and Codex.
Record every outcome.

03

Review

Inspect by task family.
Check human acceptance.
Share the evidence.

[View evidence on GitHub ↗](#)

Small exploratory samples and repeated designs cannot establish a universal ranking. Full outcomes, retries and measurement limits remain in the evidence package. Early-access results are not a GA retest.

More room to **build.**

The exciting signal: **less generated output with passing results** in selected reasoning and coding comparisons.

Next: longer workflows and more autonomous execution.
What are you seeing in your own builds?

An early-access field study of Vega-alpha and GPT-5.6 Sol. Evaluate your own acceptance criteria and current deployed versions before generalizing these selected findings.