

# The task decides the upgrade.

Where the early-access model saved tokens.

Where the advantage disappeared.

An Eliza field evaluation · Teja

Findings from early-access Vega-alpha testing against GPT-5.6 Sol. The launch version has not been retested.

# One model.

# Opposite token results.

## Exact constraint reasoning

3 designs · 6 passing pairs

## Earlier API task packets


15 designs · 15 passing pairs

## Public coding subset

16 designs · 16 passing pairs

## Staged repository work

3 clean workflows · 3 passing pairs



Same effort within each row: high vs high.

Matched passing pairs; change in summed output tokens, including reasoning. Separate task families, not a pooled model effect. Repeats are not new designs.

# API and Codex answer different questions.

## Direct API

**332** candidate workflows

Explicit prompts and settings.

Provider usage and task-specific checks.

## Codex

**80** candidate assignments

Model + app context + tools + cache.

Captured inputs differed by model.

## Vega vs Sol in Codex

Selected high-effort subsets

**16.2% less**

**coding output**

3 designs / pairs · both 3 / 3 pass

**24.2% more**

**FDE task output**

3 designs / pairs · recorded passes 2 / 3 each

Codex examples are synthetic packet tasks, not long repository work. FDE grading included model rubrics. Different tasks and runtime context prevent an API-versus-Codex causal comparison.

# Token efficiency matters after it works.

100 / 100

**Sol objective passes**

Expanded API campaign · assigned tasks

98 / 100

**Vega objective passes**

Expanded API campaign · assigned tasks

---

**Two Vega attempts hit the response cap.**

Those failures stay in the result.

200 assigned attempts across 22 selected designs and multiple effort settings. Both failures were one repository design at high / xhigh. This is not established production-quality equivalence.

# The advantage can get small.

4.1%

**fewer Vega output tokens**

One repository design · 2 pairs · high

37.9%

**longer completion time**

Both models passed 2 / 2

---

The latest follow-up showed a small token difference.

Fresh attempts used a larger budget, longer timeout and serial execution. Recovery does not prove the old cap caused the failures. Time is median end-to-end elapsed time, not decoding speed.

Measure the full  
cost of the task.

29%

**fewer Vega output tokens**

Versus Sol · long-context workflows

Here, input dominated  
the **provider cost estimate**.



~1.4%

**lower reference cost**

Same price assumptions for both

High · 2 designs · 4 passing pairs. Shared Sol-rate scenario; Vega pricing was unverified. Provider calls only: review, tool hosting and orchestration costs are excluded.

# The evidence should travel with the claim.

01

## Task

Save the prompt.  
Define acceptance.  
Record effort and limits.

02

## Run

Keep provider usage.  
Separate API and Codex.  
Count failures and retries.

03

## Review

Inspect by task family.  
Check human acceptance.  
Publish the limitations.

[Companion evidence package ↗](#)

Small exploratory samples and repeated designs cannot establish a universal model ranking. The package separates candidate outcomes, measurement boundaries and public-release limitations.

# Test the workflow. Then choose the model.

At Eliza, the decision is the workflow:  
**accepted result, elapsed time,  
and total cost including repair.**

**Route by task.**

**Keep the failures in the ledger.**

---

An early-access field study of Vega-alpha and GPT-5.6 Sol. The next decision should use your own acceptance criteria and current deployed versions.