

The task decides the upgrade.

Where the early-access model saved tokens.

Where the advantage disappeared.

364 candidate attempts · 18.4M recorded tokens

An Eliza field evaluation · Teja

Input + output, including cached context, across both models. Early-access Vega-alpha vs GPT-5.6 Sol; not a launch-version retest. Attempts are not independent designs.

One model.

Opposite token results.

Exact constraint reasoning

3 designs · 6 passing pairs

Earlier technical API packets

14 designs · 14 passing pairs

Public coding subset

16 designs · 16 passing pairs

Staged repository work

3 clean workflows · 3 passing pairs

Fewer Vega tokens

More Vega tokens

–64.2%

–1.3%

+1.1%

+40.7%



Same effort within each row: high vs high.

Matched passing pairs; change in summed output tokens, including reasoning. Separate task families, not a pooled model effect. Repeats are not new designs.

API and Codex need separate readouts.

Public comparison scope: 322 API workflows · 42 Codex assignments


WITHIN CODEX · OUTPUT TOKENS INCLUDING REASONING · VEGA VS SOL



FDE-specific lanes excluded; full historical ledger preserved. Codex used synthetic packet tasks with unequal runtime inputs. Different tasks and runtime context prevent an API-versus-Codex causal comparison.

Token efficiency matters after it works.


Each square = one assigned attempt



100 / 100

Sol objective passes

0 failed



98 / 100

Vega objective passes

2 response-cap failures



Those two failures stay in the result.

200 assigned attempts across 22 selected designs and multiple effort settings. Both failures were one repository design at high / xhigh. This is not established production-quality equivalence.

The advantage can get small.

Output tokens **4.1% less**

Elapsed time **37.9% longer**



22,529 → 21,603 tokens · sum over 2 pairs

155.6 → 214.5 seconds · median elapsed

Within each measure, Sol = 100. One repository design · 2 high-effort pairs · both models 2 / 2 pass.

A small token difference. **A longer wait.**

Fresh attempts used a larger budget, longer timeout and serial execution. Recovery does not prove the old cap caused the failures. Time is median end-to-end elapsed time, not decoding speed.


Measure the full cost of the task.

29%

fewer Vega output tokens

~1.4%

lower reference cost



Input dominated the provider cost estimate.

High · 2 designs · 4 passing pairs. Shared Sol-rate scenario; Vega pricing was unverified. Provider calls only: review, tool hosting and orchestration costs are excluded.

The evidence should travel with the claim.

01

Task

Save the prompt.
Define acceptance.
Record effort and limits.

02

Run

Keep provider usage.
Separate API and Codex.
Count failures and retries.

03

Review

Inspect by task family.
Check human acceptance.
Publish the limitations.

[View evidence on GitHub ↗](#)

Small exploratory samples and repeated designs cannot establish a universal model ranking. The package separates candidate outcomes, measurement boundaries and public-release limitations.

Test the workflow. Then choose the model.

At Eliza, the decision is the workflow:
**accepted result, elapsed time,
and total cost including repair.**

Route by task.

Keep the failures in the ledger.

An early-access field study of Vega-alpha and GPT-5.6 Sol. The next decision should use your own acceptance criteria and current deployed versions.