

Where Astra improves on Sol.

Selected early-access results from our
API and Codex evaluations at Eliza.


364 candidate attempts · 18.4M recorded tokens

Early access, real experiments · Teja

Recorded tokens include input, output and cached/repeated context across both models. Vega-alpha is the early-access model; results vary by task.

Exact reasoning.

64.2% fewer output tokens.



Compared with Sol. **The same acceptance checks passed.**



Summed output includes reasoning. Small exploratory sample; the same effort setting does not imply equal compute.

Less output than Sol.

In Codex, too.

Public comparison scope: 322 API attempts · 42 Codex assignments


ASTRA VS SOL · OUTPUT TOKENS, INCLUDING REASONING



Synthetic packet tasks; runtime context differed by model. Read these results separately from the API tests.

Long context.

29.0% fewer output tokens.



Both models passed across the two context sizes.

Nominal 96k and 192k context sizes. Output includes reasoning; input and cached context are excluded from this comparison.

A more natural delegation workflow.

Eliza teammates described Astra using
subagents more naturally and needing
less steering on longer implementations.

Plan → Delegate → Coordinate

A practitioner signal worth testing next.

Qualitative Eliza engineering feedback. Autonomous delegation and coordination were not measured in the benchmark.

Match the model to the task.

Exact constraint reasoning

3 designs · 6 passing pairs

Earlier technical API packets

14 designs · 14 passing pairs

Public coding subset

16 designs · 16 passing pairs

Staged repository work

3 clean workflows · 3 passing pairs

Fewer Astra tokens

More Astra tokens



Same effort within each row: high vs high.

Separate task families and matched passing pairs. Output includes reasoning. Gains vary by workload; repeats are not new designs.

The evidence should travel with the claim.

01

Task

Save the prompt.
Define acceptance.
Record effort and limits.

[View evidence on GitHub ↗](#)

02

Run

Keep provider usage.
Separate API and Codex.
Record every outcome.

03

Review

Inspect by task family.
Check human acceptance.
Share the evidence.

Full outcomes, retries and limits are linked. Small exploratory samples and repeated runs do not establish a universal ranking.

A promising upgrade. For the right tasks.

On selected tasks, Astra used
fewer output tokens than Sol
with the same checks passed.

Next: longer workflows and more autonomous execution.
What are you seeing in your own builds?

Validate the current models on your own tasks and acceptance criteria before generalizing these early-access findings.