Where Astra improves on Sol.

Selected early-access results from our API and Codex evaluations at Eliza.

18.4M recorded tokens across both models

Early access, real experiments · Teja

Recorded tokens include input, output and cached/repeated context across both models. Results vary by task.

Exact reasoning.
64.2% fewer output tokens.

Summed output tokens, including reasoning

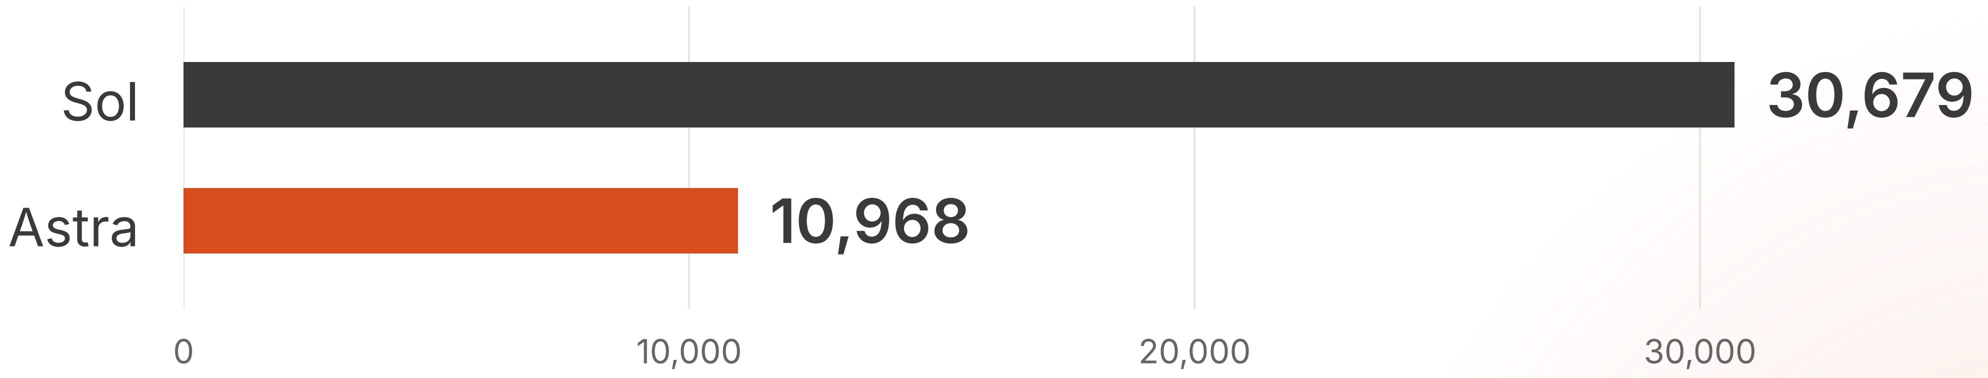

High reasoning effort · Matched tasks and acceptance checks

Compared with Sol. **The same acceptance checks passed.**

Summed output includes reasoning. Small exploratory sample; the same effort setting does not imply equal compute.

API AND CODEX · SEPARATE EVIDENCE SURFACES

Less output than Sol.
In Codex, too.

ASTRA VS SOL · OUTPUT TOKENS, INCLUDING REASONING

Coding 16.2% less

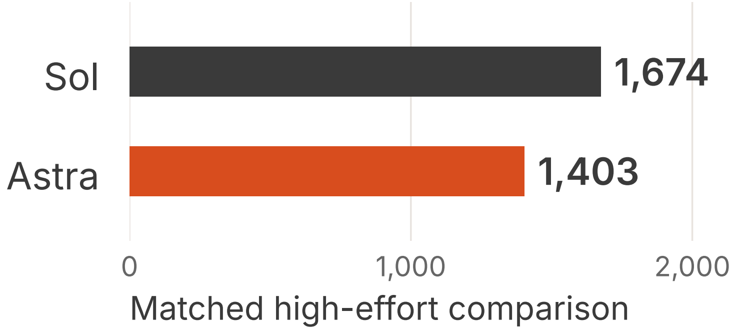

Reasoning 25.8% less

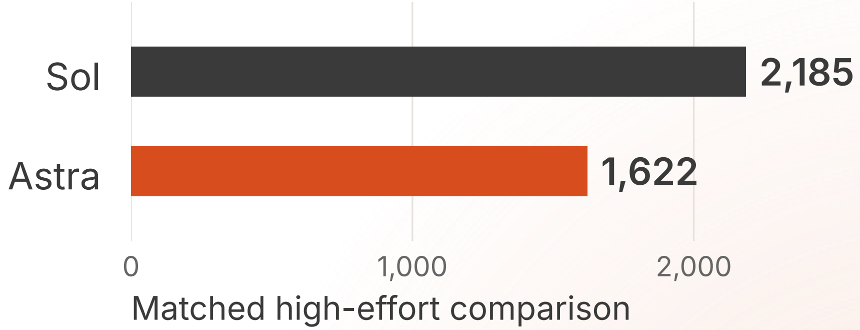

Synthetic packet tasks; runtime context differed by model. Read these results separately from the API tests.

Long context.
29.0% fewer output tokens.

Summed output tokens, including reasoning

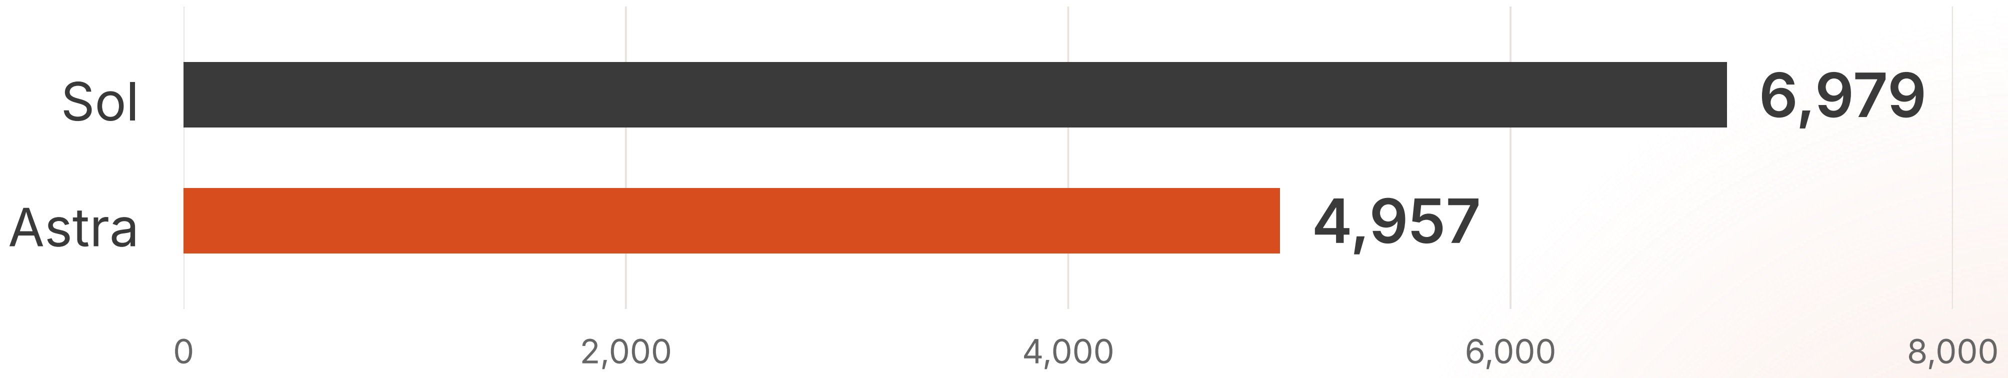

High reasoning effort · Matched tasks and acceptance checks

Longer context. **The same acceptance checks passed.**

Output includes reasoning; input and cached context are excluded. Test details and context sizes are linked in the evidence.

A more natural delegation workflow.

Eliza teammates described Astra using **subagents more naturally** and needing **less steering** on longer implementations.

Plan → Delegate → Coordinate

A practitioner signal worth testing next.

Qualitative Eliza engineering feedback. Autonomous delegation and coordination were not measured in the benchmark.

Match the model
to the task.

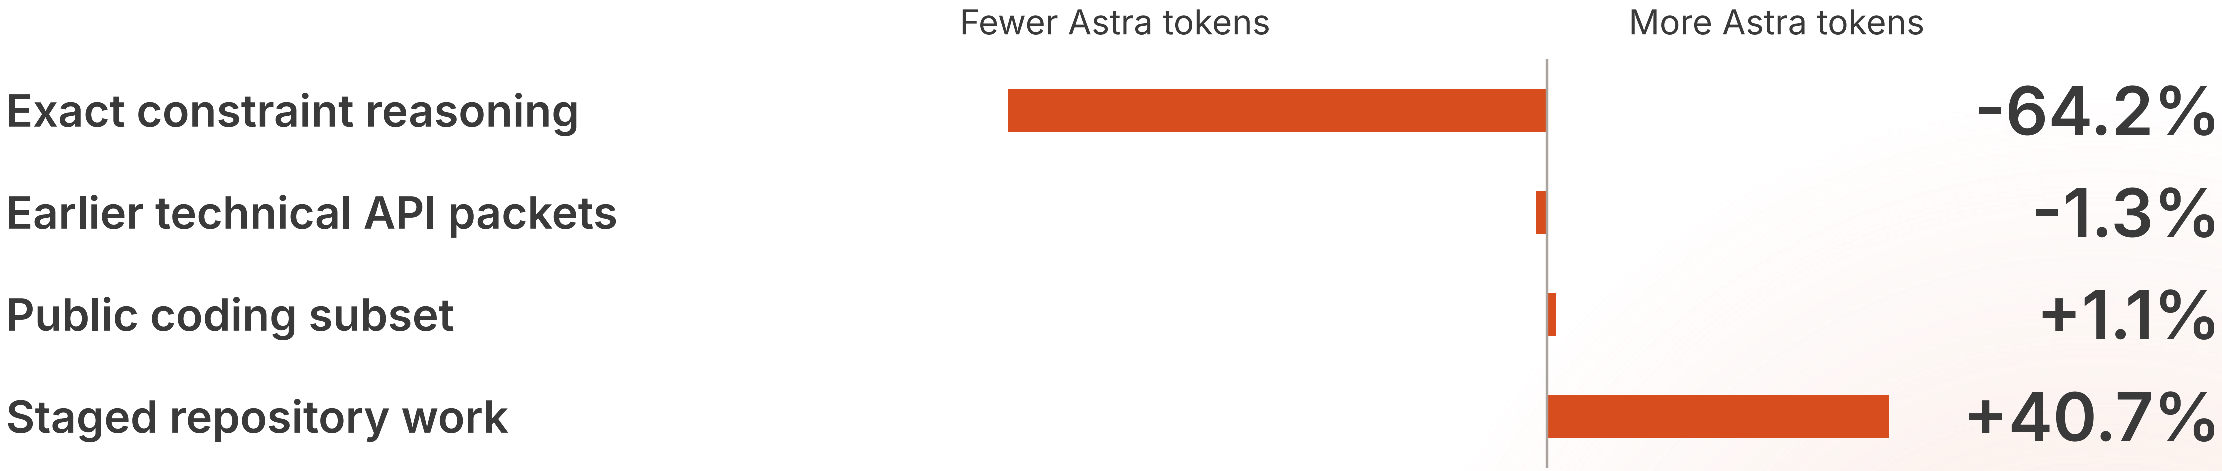

Same effort within each row: **high vs high.**

Separate task families and matched passing pairs. Output includes reasoning. Gains vary by workload; repeats are not new designs.

AN INSPECTABLE EVALUATION

The evidence should travel with the claim.

01

Task

Save the prompt.

Define acceptance.

Record effort and limits.

02

Run

Keep provider usage.

Separate API and Codex.

Record every outcome.

03

Review

Inspect by task family.

Check human acceptance.

Share the evidence.

Methodology + data + reproducible checks

Full outcomes, retries and limits are included in the companion evidence package. Small exploratory samples and repeated runs do not establish a universal ranking.

A promising upgrade.
For the right tasks.

On selected tasks, Astra used **fewer output tokens than Sol with the same checks passed.**

Next: longer workflows and more autonomous execution.
What are you seeing in your own builds?

Validate the current models on your own tasks and acceptance criteria before generalizing these early-access findings.